

Public contracting for private innovation: Government capabilities, decision rights, and performance outcomes

Joshua R. Bruce¹  | John M. de Figueiredo² | Brian S. Silverman³

¹Department of Sociology, Duke University, Durham, North Carolina

²School of Law and Fuqua School of Business, Duke University, Durham, North Carolina

³Rotman School of Management, University of Toronto, Toronto, Ontario, Canada

Correspondence

Joshua R. Bruce, Department of Sociology, Duke University, 417 Chapel Drive, Box 90088, Durham, NC 27708.

Email: joshua.bruce@duke.edu

Funding information

National Science Foundation, Grant/Award Number: 1061600 1443014

Research Abstract: We examine how the US Federal Government governs R&D contracts with private-sector firms. The government chooses between two contractual forms: grants and cooperative agreements. The latter provides the government substantially greater discretion over, and monitoring of, project progress. Using novel data on R&D contracts and on the technical expertise available in specific government bureau locations, we test implications from the organizational economics and capabilities literatures. We find that cooperative agreements are more likely to be used for early-stage projects and those for which local government scientific personnel have relevant technical expertise; in turn, cooperative agreements yield greater innovative output as measured by patents, controlling for endogeneity of contract form. The results are consistent with multitask agency and transaction-cost approaches that emphasize decision rights and monitoring.

Managerial Abstract: When one private firm outsources an R&D project to another, it can use a range of sophisticated contractual provisions to elicit proper innovative effort. However, government entities are often constrained from employing such provisions due to legal and regulatory restrictions. Policymakers thus face a difficult challenge when contracting with private firms for innovation. We study the US Federal government's R&D contracts, which are restricted to two contractual types: "grants," which offer little in-process oversight, and "cooperative agreements," which provide decision rights during the project. We demonstrate that policymakers can enhance outcomes by using cooperative agreements for earlier-stage, higher-uncertainty projects, but only when government scientists with relevant expertise are located near the firm's R&D site.

KEY WORDS

capabilities, contracts, innovation, public-private, R&D

1 | INTRODUCTION

How can public-sector actors engage in value-creating arrangements with private-sector entities? Strategic management scholars have illuminated a range of benefits that can arise from collaboration between private firms, and numerous solutions for overcoming the challenges of managing such collaboration. In principle, public–private contracting should be amenable to many of the same solutions. However, in practice, contracting that involves the public sector is characterized by strong rigidity (Moszoro, Spiller, & Stolorz, 2016), due to both political pressure from third parties (Kivleniece & Quelin, 2012; Spiller & Moszoro, 2014) and more general bureaucratic structures (Moe, 1990). Consequently, although theories prevalent in strategic management can inform public–private contracting, the distinct characteristics of the public sector—which preclude the use of several creative contracting mechanisms implicated in such theories—highlight the need for nuanced extension rather than wholesale application of these approaches.

This paper extends contracting theory to the public–private sphere in a specific setting: public contracting for private innovation. Government entities frequently encourage the development of specific innovations that are deemed necessary to achieve public aims. For example, in 2016 the US federal government funded \$148 billion in R&D (Hourihan & Parkes, 2015), at least \$30 billion of which was devoted to contracts for private innovation regarding specific goals such as preclinical trials for new pharmaceutical drugs, studies of chemical toxicity in water, and the development of systems on the Mars rover missions.¹

Contracting for innovation faces a central problem: when a client organization pays a research firm to pursue a specific project, there is a risk that the research firm will divert the payment to pursue its own interests. If the research project fails, it is difficult for the client organization to distinguish between research-firm malfeasance and bad luck. This problem is exacerbated as the uncertainty surrounding the project increases. Organizational economists offer several prescriptions to overcome this problem. However, virtually all of the proposed solutions depend on substantial flexibility in contract design: judicious allocation of property rights (Aghion & Tirole, 1994), menus of fixed fees and royalty payments (Hegde, 2014), appropriate investments in equity (Oxley, 1997), or sophisticated contractual provisions to deal with a range of contingencies that might arise (Reuer & Ariño, 2007). These solutions are often precluded by the rigidities of public-sector contracting, such as the need for standardized contracts, the desire to prevent officeholders from self-dealing, the requirement for procedural adherence and transparency by agencies and civil servants, the presence of public sector unions, and the attempt to insulate the bureaucracy from the political pressures of third parties (Moszoro et al., 2016). Given these constraints, how can public–private contracts for innovation overcome contractual hazards in order to create value?

¹Similarly, in 2012 the United Kingdom allocated £1.6B (18%) of its £8.7B government R&D budget to private firms (National Audit Office, 2013).

We invoke and extend a solution that stems from recent analyses of privatization: government retention of specific decision rights even while the private actor retains residual control rights. Specifically, Hart, Shleifer, and Vishny's (1997) influential model of privatization predicts that privatized services, for which the private provider owns residual control rights including decision rights, will have lower costs but also lower quality than publicly provided services. Yet Williamson (1999) proposes that overlapping decision-making authority between the private entity and public bureau may ameliorate quality degradation. Empirically, Cabral, Lazzarini, and de Azevedo (2010, 2013) find that privatized prisons in Brazil exhibit quality equivalent to that of publicly run counterparts, attributing this to the presence at each private prison of a government supervisor, or "warden," whose job is to review and occasionally overrule decisions about those aspects of operations that could adversely affect quality. This preserves quality as long as the government warden is motivated to make good decisions.

We extend this logic to make three predictions about public contracting for private innovation. As contracting hazards increase due to project uncertainty, the government agency is more likely to include an information exchange and decision-rights mechanism in the contract. However, this mechanism will only be implemented if the government agency has employees with the requisite subject-matter expertise to make effective and appropriate project decisions. Finally, conditional on project uncertainty and government expertise, agreements with such decision rights will be more successful at generating patented innovations (which is an explicit goal of such projects) because of more effective monitoring of private-sector R&D efforts by the government.

We test these predictions using a sample of more than 4,000 R&D contracts between US federal agencies and private firms. Similar to many countries, US law generally restricts these contracts to take one of two forms: a "grant," which affords the government no in-process decision rights, and a "cooperative agreement," in which government employees have substantial in-process decision rights. Using novel data on the technical expertise of government agency personnel located in geographic proximity to the private firm's R&D site, we find support for our theoretical predictions. Notably, (a) earlier-stage projects (which are likely to entail greater uncertainty) are more likely to be governed by cooperative agreements than by grants, (b) agencies rely on cooperative agreements more readily when they have relevant technical capabilities near the R&D site, (c) cooperative agreements perform better than grants in terms of patents generated and the citations to these patents, and (d) although cooperative agreements perform better than grants overall, those projects that were governed by grants would not have been as productive as current cooperative agreements had they been organized as cooperative agreements. We then consider a number of alternative theoretical explanations, econometric specifications, and data measurements, and find the results are robust to these approaches.

This study makes three contributions to the theoretical literature. First, while most analyses of privatized services have focused on the in-house versus privatization decision, we extend the logic to consider variations in privatized governance based on different characteristics of projects. Second, while recent literature on hybrid governance highlights the condition that the government monitor must be *willing* to monitor effectively, we propose that, for projects that rely on highly idiosyncratic knowledge, the monitor must be both *willing and able*—that is, must have the requisite skills to monitor and make good decisions. Third, this focus on requisite skills extends the literature on government capabilities: whereas prior research on value creation in public-private collaboration has tended to emphasize a government entity's capabilities in contracting (e.g., Klein, Mahoney, McGahan, & Pitelis, 2013), we extend this to consider how a government entity's technological capabilities

influence the organization of privatized services.² Through these contributions, we further flesh out the implications of Moszoro et al.'s (2016) insights about rigidity in public–private contracting.

The paper also makes an empirical contribution. We develop measures of government innovative capabilities, based upon the specific technical expertise of government scientists, engineers, doctors, and researchers as measured by 9.5 million person-year observations of government personnel data. We then geo-locate those capabilities throughout the United States, essentially creating a map of government capabilities along 59 expertise dimensions. To the best of our knowledge, this is the first time that government capabilities have been measured in such a microanalytic way. We explore in the conclusion some ways this measure can be used to further inform the innovation and strategic management literatures.

The paper proceeds as follows. In Section 2, we analyze the public-sector challenge of contracting effectively for R&D, ultimately generating predictions regarding the use of government decision rights in contracts for innovation. Section 3 provides institutional detail on the US empirical setting. Section 4 introduces our data, model, and empirical strategy. Section 5 presents empirical results, and Section 6 offers a brief discussion and conclusion.

2 | CONTRACTING FOR RESEARCH—THE PRIVATE–PRIVATE VERSUS PUBLIC–PRIVATE CONTEXT

2.1 | The private–private context

The market for technology suffers from several well-documented defects (Arora & Gambardella, 2010). The R&D process is commonly characterized by several features that create contractual hazards, including uncertainty, noncontractible effort, tacit knowledge, and appropriability concerns. Organizational economics theories generally agree that contracting difficulties rise monotonically with these characteristics. Given such difficulties, scholars have devoted substantial attention to explicating contractual mechanisms that private entities can use to efficiently govern R&D transactions.

Consider an example in which a pharmaceutical firm seeks to contract with a biotechnology firm for R&D into a new drug. The client firm pays the research firm to conduct a set of specified research tasks. But it is nearly impossible for the client firm to observe the effort that the research firm's employees actually devote to the tasks. R&D is an uncertain endeavor, so if the research firm does not generate the desired innovation, it is difficult to tell whether this was the result of insufficient effort or bad luck. Even when the innovation is developed, if transfer to the client firm requires the provision of attendant tacit knowledge, then it is difficult to monitor whether the researchers are making a good-faith effort to provide this knowledge (Hegde, 2014), especially if either firm has latent concerns that proprietary knowledge outside the scope of the contract will "leak" to the other party during the course of the endeavor (Oxley & Wada, 2009).

Given these challenges, how might the firms successfully govern their exchange? One prescription is to judiciously assign property rights so as to elicit noncontractible effort as effectively as is feasible (Aghion & Tirole, 1994; Grossman & Hart, 1986). Lerner and Merges (1998) test this empirically by exploring the pattern of property-rights assignment in R&D contracts between pharmaceutical companies and biotechnology firms. They find modest evidence that these contracts do indeed

²Relatedly, Decarolis, Giuffrida, Iossa, Mollisi, and Spagnolo (2018) examine public–private procurement contracting but focus on the competencies of private firms rather than the public sector.

assign more property rights to the biotech firm when projects are earlier-stage (and hence are more uncertain and require more noncontractible effort from the biotech firm). Lerner and Malmendier (2010) consider termination options that distribute rights to R&D results when projects characterized by unobservable effort also generate observable milestones, finding that such termination options also appear more frequently in contracts for earlier-stage projects than in contracts for later-stage projects.³

An alternative prescription is to implement a combination of fixed fees and royalty payments to align the firms' incentives. Since royalties depend on successful commercialization of an innovation, they can provide a strong incentive to the research firm to both conduct the requisite R&D and devote effort to transferring the results to the client firm (Xiao & Xu, 2012). Although reliance on royalties shifts risk to the research firm, which in many models is more risk-averse than the client firm, the benefits of incentive alignment outweigh the attendant costs for sufficiently high levels of uncertainty and noncontractible effort. In a study of biomedical invention, Hegde (2014) finds systematic patterns of complex royalty payments between commercializing firms and inventors that are consistent with theoretical predictions.

A third prescription is to judiciously use equity investments to align incentives, direct effort, and protect knowledge (Pisano, 1990; Teece, 1986). While nonequity arrangements such as licensing contracts will suffice for high-appropriability or low-tacit-knowledge research in the presence of low uncertainty, equity joint ventures will be used to govern research agreements with higher levels of contractual hazards (Oxley & Wada, 2009). Shared ownership of the collaborative venture implies shared ownership of the attendant profits, thus aligning the firms' incentives regarding success of the venture. Equity arrangements also provide formal monitoring and, in particular, decision-making authority over the research effort (Reuer, Ariño, & Mellewigt, 2006). These predictions have been borne out in specific industry settings (Sampson, 2004a) and multi-industry studies (Oxley, 1997). Going beyond the governance-choice decision, Sampson (2004b) also finds that R&D alliances that are organized according to transaction-cost precepts generate more patented innovations than those that are organized inappropriately.

Finally, complex contractual provisions may be employed to coordinate and control effort in the face of uncertainty. Contractual features such as contingency payments can elicit effort and align incentives in contractual relationships, while provisions that specify responses to potential contingencies can restrict opportunistic behavior (Argyres, Bercovitz, & Mayer, 2007). Contractual clauses that thus effectively address hazards can dramatically increase contractual effectiveness (Anderson & Dekker, 2005) and facilitate resolution of disagreements (Lumineau & Malhotra, 2011). Alternatively, judicious assignment of decision rights and monitoring provisions can dramatically influence the effectiveness of incentives and the performance of the project (Arruñada, Garicano, & Vázquez, 2001; Athey & Roberts, 2001; Reuer & Devarakonda, 2016).

In general, then, contracting between private firms is frequently facilitated by a range of governance mechanisms including judicious allocation of property rights, complex royalty schemes, equity holdings, and/or sophisticated contingent contracts with attendant decision rights. These mechanisms support a substantial market for technology both within and across nations (Arora & Gambardea, 2010).

³A unilateral termination option for the client firm encourages the research firm to devote appropriate effort to the project, while a termination fee set at an appropriate level discourages the client firm from strategically terminating the project. This option can align incentives between client firm and research firm.

2.2 | The public–private context

At first glance, one might expect that the above prescriptions are straightforwardly applicable to public–private contracting. However, in many countries, strict rules and processes constrain the form of public contracts for innovation. In the United States, as in several other Organization for Economic Cooperation and Development (OECD) countries, government entities are prohibited from owning property rights in the resulting innovations, paying royalties to the contracted firms, or taking equity in these firms. Strict contracting policies also hinder attempts to craft project-specific contractual provisions.⁴ Thus, the most common levers available to private–private contracts for research are unavailable in public–private research contracts. These constraints reflect the stylized fact that public-sector contracts tend to be far more rigid than their private-sector counterparts (Moszoro et al., 2016). This enduring feature of public bureaucracy (Boyne, 2002) is often attributed to a desire to restrict public agents' ability to engage in self-dealing (Lan & Rainey, 1992); or, alternatively, to concerns about political pressure from third parties (Spiller & Moszoro, 2014). This then leads to processes in government which are procedurally onerous and substantively transparent, often leading to inefficiency in the government by design (Moe, 1989). Because of these substantial procedural requirements and limited resources for government contracting, the government tends to favor standardized, rather than customized, contracts for many purposes, limiting the ability of the government to employ specialized terms (Miller, 1955).

This rigidity is manifest in Hart et al.'s (1997) incomplete-contract model of privatization. In this model, a government actor chooses between delivering a service through in-house provision or through a contract with a private provider. The service requires investment in an asset and then operation using that asset. Of particular relevance, property rights over the asset cannot be divided, but rest entirely with either the government or the firm, and payment to the firm is limited to a fixed fee that can be renegotiated upward if the quality of the service is increased. Given these blunt levers, the private firm has a strong incentive to lower the cost of provision, even at the expense of quality, while an in-house provider has little incentive to improve either cost or quality.⁵ Consequently, Hart et al. predict that privatized services will have lower costs but also lower quality than their publicly provided counterparts. Levin and Tadelis (2010) find that municipalities are less likely to outsource services for which quality is important yet noncontractible, concluding that this is consistent with the Hart et al. (1997) model. This provides a pessimistic assessment of the feasibility of public–private contracting for innovation, given its reliance on noncontractible effort.

Yet Cabral et al. (2010, 2013), building on Williamson's (1999) rejoinder to Hart et al. (1997), find that privatized prisons in two Brazilian states exhibit quality that is equal to or better than that of their publicly run counterparts, even while enjoying lower costs. They propose that the key quality-protection mechanism is the appointment to each private prison of a government “warden” whose job is to monitor prison operation and ensure that it adheres to specified minimum quality standards. As long as the warden remains committed to her task—that is, she is not bribed by the private firm—then this “hybrid” form of private operation and public supervision appears to solve the problem of quality deterioration.⁶

⁴For example, similar to the United States, Canada's federal contracts for R&D take the form of either grants or “contributions,” which are largely analogous to US cooperative agreements.

⁵The private firm reaps the entire benefit from cost reduction, but only incurs a fraction of the benefit to quality improvement because it must bargain with the government ex post for fee increases associated with improved quality. The private firm thus “overinvests” in cost reduction, yielding a socially suboptimal level of quality.

⁶Although the Brazilian prison setting only allows comparison of public to private-hybrid prisons, it should be the case that a purely private prison would have lower costs than the private-hybrid prison. Some of this would be due to quality-shading efforts that are socially destructive. But some should be due to lower effort to invest in cost-reduction by the private-hybrid, given that the warden may sometimes erroneously negate a valid cost-reduction scheme.

To the extent that public contracting for innovation is characterized by the constraints embodied in the Hart et al. (1997) model, perhaps the sole available lever is the prospect of government supervision of the research project, analogous to the government prison warden. Yet one difference stands out in the innovation setting. Cabral et al. (2010, 2013) implicitly assume that the government warden understands the causal mechanisms linking cost-reduction and quality shading—in essence, she knows which actions by the private agent are good and which are bad. However, evidence indicates that an organization's possession of relevant technological capability is necessary for it to appraise the value of external research (Cassiman & Veugelers, 2006): “[t]he ability to evaluate...outside knowledge is largely a function of the level of prior related knowledge” (Cohen & Levinthal, 1990, p. 128).

Indeed, precisely for this reason, an organization's technological capability in a particular sphere can influence its competence at contracting in that sphere. Mayer and Salomon (2006) study an IT firm's decisions to complete client projects with in-house or outsourced teams, finding evidence that strong technological capability in, for example, mainframe technology allows the firm to outsource on mainframe-related projects in the face of contractual hazards. They conclude that the firm is better able to manage an external contract when it has technological capabilities that enable it to anticipate problems and monitor outcomes. Building on this idea, Argyres and Mayer (2007) propose that the technological (i.e., engineering) expertise of a firm's employees is particularly relevant to establishing effective interfirm communication flows in research contracts. Thus, in the context of public contracting for innovation, the government supervisor must have the requisite technological expertise to know which actions are good and bad—in other words, if the supervisor is willing but not able to make good decisions, then public supervision is a hindrance.⁷

In sum, extending the predictions concerning decision rights and monitoring above, we expect to find two patterns in contract choice: research projects that are more uncertain in outcome are more likely to be governed by contracts that afford greater public supervision (i.e., cooperative agreements), as are research projects for which the available public personnel have relevant expertise. We further expect research projects governed according to the above precepts will outperform those that are not in terms of innovative output.

3 | GOVERNMENT CONTRACTING FOR RESEARCH: INSTITUTIONAL DETAILS

The US federal government is composed of 381 agencies, which in turn are composed of 874 bureaus. Although formally overseen by the Executive Branch of the government, agencies pursue their own research and development agendas, each determined by a variety of different considerations. To meet their required objectives, bureaus often determine that specific research endeavors would require expertise beyond that available within the federal government.⁸ In such cases, the bureau contracts with outside entities for the requisite research effort.

⁷There have been recent calls to incorporate capabilities-related insights more centrally in research on public-private interactions (e.g., Klein et al., 2013; Quelin, Kivleniece, & Lazzarini, 2017), as well as explicit efforts to introduce contracting capabilities to such research (e.g., Cabral, 2017). This study's emphasis on government's technological capabilities expands the range of public capabilities that are considered relevant to public-private contracting.

⁸McKenna (2006, pp. 103-105) describes the government's strategic decision, at the beginning of the US space program, to rely on external expertise rather than try to employ all necessary experts within NASA.

The process begins when a bureau's program office issues a call for research proposals, or CFP (see Supporting Information Appendix S1: The Grant-Making Process for a detailed description of the CFP process). The CFP outlines the motivation for the research project, the statutory authority for the agency to conduct the research, a list of requirements, milestones, expectations, and objectives of the project, a list of eligibility requirements for the private contractor, and a description of how the project will be managed. The CFP may specify that the research will be conducted through a grant, a cooperative agreement, or either. As Supporting Information Appendix S1 describes, the CFP process is virtually identical across the two types of contract. In both cases, all property rights resulting from the contracted research are owned by the contracted entity while the US government receives a royalty-free license. This precludes the judicious allocation of property rights and the use of royalty schemes to elicit effort.⁹ However, for the purposes of this study, there are three key differences between the governance forms.

The first difference is the degree of cooperative effort between the government agency and private firm. As stipulated in the Federal Grant and Cooperative Agreement Act of 1977 (FGCAA) and the Code for Federal Regulation (CFR), grants do not provide for "substantial involvement" between government employees and the firm, whereas cooperative agreements do.¹⁰ This is reinforced by each agency's own guidance documents. For example, Section 3 of the NASA Grant and Cooperative Agreement Manual (2016, p. 3) notes that unlike a grant, a cooperative agreement should be used if "substantial involvement is expected between the executive agency and the...other recipient when carrying out the activity contemplated in the agreement."¹¹

The second difference relates to the disparate pattern of decision rights assigned to private firm and government. Grants typically allow the recipient firm's principal investigator to make virtually all key decisions during the research project, subject to compliance with federal regulations. In cooperative agreements, decision rights are more evenly distributed between government and firm personnel. Daily decisions are often jointly determined by both parties. For example, in a cooperative agreement between the National Cancer Institute (NCI) and GlobeImmune, Inc. to develop yeast-based tarmogens for cancer immunotherapy, the NCI and GlobeImmune each had its own Principal Investigator. This role, as specified in the agreement, was to be "person(s) designated by the Parties who will be responsible for the scientific conduct of the Research Plan."¹²

Cooperative agreements also often provide the government with the right to terminate a project before its official completion should the government's principal decision-maker on the project determine that its progress is not satisfactory. Indeed, the Department of Energy's Model Cooperative Agreement in Energy Efficiency and Renewable Energy contains not only regular review meetings for the government, but also contains a section that grants the government "go/no-go decisions" and decision-making authority at key milestones in the project.¹³ In contrast, although the government

⁹Some federal procurement agreements also involve R&D effort by the private vendor. In these agreements, called "contracts," the federal government funds the vendor's R&D as "work for hire" and receives ownership of any resulting patents. We exclude these from this study for two reasons: they are not designed to support significant R&D; and their different (although still rigid) allocation of property rights would conflate the incentives affecting contract performance.

¹⁰See "Implementation of the Federal Grants and Cooperative Agreements Act of 1977, Office of Management Budget, August 18, 1978, *Federal Register* 43(161): 36860–36,865." "Substantial involvement" does not have a formal regulatory definition, but it is described variously as entailing direction and redirection of the technical aspects of the project as a whole; sharing responsibility with the firm for the management, control, direction, and performance of the project.

¹¹NASA *Grant and Cooperative Agreement Manual*, Revised September 16, 2016.

¹²"Preclinical and Clinical Development of GlobeImmune, Inc's Proprietary Yeast-Based Tarmogens Expressing Tumor-Associated Antigens for Cancer Immunotherapy," between the National Cancer Institute, NIH, and GlobeImmune, Inc., signed May 8, 2008.

¹³"Model Cooperative Agreement," Contractual Term 7D. US Department of Energy, Energy Efficiency and Renewable Energy Program, February 19, 2013.

can in principle decide to withhold subsequent funding payments from an in-process grant, this tactic is cumbersome to implement and rarely employed.

The third substantive difference between grants and cooperative agreements stems from the decision-rights difference, and relates to the degree of information that passes between the private firm and the government. Although a grant is awarded to a recipient firm through a rigorous review process, during the project the recipient is only required to provide the agency with periodic (often annual) reports of progress made on the grant's objectives. After the project is completed, the recipient has a finite amount of time (usually 90 days) to file a final report of accomplishments.

In cooperative agreements, the private firm is expected to provide information to the government on a much more frequent basis. Given that decisions regarding project tasks are sometimes made as frequently as daily, information must flow almost continuously to support informed government decision-making. In those cases where government and private firm scientists work closely together, this can occur informally through the collaborative effort. In cases where this collaboration does not occur consistently, agreements stipulate formal obligations to provide for communication. Thus, in contrast to research grants where researchers provide information to the government at specified, infrequent intervals, cooperative agreements stipulate more rapid communication and flow of information.

For example, a cooperative agreement between the Department of Energy and Mascoma Corporation, for a project to demonstrate feasibility of biorefining technology using plant biomass, specified that "in order to adequately monitor project progress and provide technical direction to the Recipient, DOE must [attend Mascoma Corporation] meetings, reviews and tests." Presumably to protect against malfeasance, the cooperative agreement further noted, "[Mascoma Corporation] shall notify the DOE Project Officer of meetings, reviews, and tests in sufficient time to permit DOE participation and provide all appropriate documentation for DOE review."¹⁴

Overall, then, research grants and cooperative agreements represent discrete structural alternatives for R&D contracting between the US federal government and private firms. Grants largely reflect canonical arms-length contracting, with little interaction during the research project except for intermittent progress reports and with the research firm retaining almost complete discretion over its allocation of effort. Cooperative agreements reflect contracting of the type prescribed above to effectively manage contractual hazards, with the government holding substantial monitoring authority and discretion over effort allocation and with the requisite information flow and interaction between the parties.

4 | DATA, MEASURES, AND EMPIRICAL STRATEGY

4.1 | Data sources

To empirically test the predictions in this study, we employ data on the characteristics of US federal government research grants and cooperative agreements, characteristics of the government bureau soliciting the project (notably the degree of relevant expertise in the local bureau offices), characteristics of the firm performing the project, and measures of innovative outcomes. We obtain these data from three sources.

¹⁴ "Demonstration of Biorefinery Application," between the Mascoma Corporation and the Department of Energy, signed September 30, 2008.

The first data set contains information on the characteristics of agreements from USASpending.gov. We downloaded all government grants and cooperative agreements (termed “assistance” in the USASpending.gov nomenclature) executed between fiscal years 2000 and 2011. Each record contains information on the governance mechanism (grant or cooperative agreement), the organization that received the funding, the principal location in which the organization would perform the research (e.g., Cincinnati, Ohio), the agency or bureau of the government that made the award (e.g., National Institute of Standards and Technology), the title and short description of the project, and other details.¹⁵ Of particular relevance, project descriptions list a set of activities necessary for the research project. We use records for only those organizations categorized as businesses in the government agreement records.¹⁶ We also remove cases where the funding agency was part of the Department of Defense or military due to data limitations.

We employ a second dataset of granted US patents, provided by PatentsView.org, to measure patent generation. We downloaded all US patents with a “government interest” indicated in the patent application. Per US regulations, patents that have any affiliation with a government unit—including any funding from that unit—must include a government-interest statement that acknowledges this affiliation. Government-interest statements refer to affiliated grants and cooperative agreements by their unique “funding identification numbers.” We then conduct an exact match against the data in the USASpending database using the funding identification numbers included in patent applications, thus identifying all patents that stipulated an affiliation with any of the contracts in the sample. We ultimately identified 1,544 patents. Of the 4,074 contracts in our sample, 508 (12.47%) led to at least one patent; 56 of these agreements supported five patents or more. Separately, to create control variables as discussed below, we use the PatentsView data to construct counts of aggregate patenting per year by each of the 383 private firms involved in any of the sample contracts.

To identify the level of relevant expertise available in specific government bureau offices, we rely on a third and relatively novel database from the US Office of Personnel Management’s (OPM) Central Personnel Data Files (CPDF). These records contain annual, individual-level information on nearly all US civil servants during the sample period (using anonymous identifier codes), including information on work location, job title and occupation, and research-related job functions.¹⁷ The CPDF allows us to measure the precise number of government employees in each office of each federal bureau who perform specific jobs. The CPDF categorizes over 800 occupations into 59 occupational groups/families, as detailed in OPM’s *Handbook of Occupational Groups and Families*.¹⁸ For example, the ‘Medical, Hospital, Dental, and Public Health Group’ includes physician assistants, nurses, nurse assistants, and doctors of dentistry, medicine, and osteopathy, among other related job titles. From these data, we compute the number of personnel in each of the 59 occupational categories at every known federal work location in the United States, geocoding each employee’s latitude and longitude. This occupation-location data is further disaggregated to the bureau level (e.g., the NIH is a bureau of the Department of Health and Human Services). Thus, we are able to identify how many of the NIH’s employees in a particular location are in the Medical, Hospital, Dental, and Public Health Group occupational category.

¹⁵A codebook for the federal assistance data set is available at <https://goo.gl/TW7QHY> (last accessed February 17, 2017). To fill in some missing project descriptions, we searched the Federal Procurement Data System and National Institutes of Health RePORTER system for federal award IDs matched to government-supported patents.

¹⁶To avoid any university affiliates misclassified as businesses, we also exclude records where the word “university” appears in the organization name.

¹⁷This data set excludes the military, US Post Office, and “sensitive” agencies (colloquially known as “three-letter agencies”) and occupations (such as US Marshals). More than 70% of US federal employees work outside the greater-DC metro area.

¹⁸Available at: <https://www.opm.gov/fedclass/GShbkocc.pdf> (last accessed February 14, 2017)

Further, for each government employee who is involved in research-related activities, broadly defined, the CPDF also includes a “functional research” category, where the set of categories includes research, development, testing and evaluation, construction, production, installation, data collection, project management, and teaching. These classifications are created by the National Science Foundation for OPM to describe the work that comprises the majority of each research employee’s time. The “research” function, for example, emphasizes early-stage research—“systematic, critical, intensive investigation directed toward the development of new or fuller scientific knowledge of the subject studied”—whereas the “testing and evaluation” function emphasizes later stage “testing of equipment, materials, devices, components, systems and methodologies under controlled conditions and the systematic evaluation of test data to determine the degree of compliance of the test item with pre-determined criteria and requirements.”¹⁹ We are thus able to identify how many of the NIH’s Medical, Hospital, Dental, and Public Health Group personnel in any location are dedicated to early-stage research, how many dedicated to development, and so on.

With data from these three sources—USASpending, PatentsView, and the CPDF database—we construct our variables.

4.2 | Variables

Our first two predictions relate to the choice of governance for a research contract. The dependent variable for testing these predictions is $Coop\ Agreement_j$, which is a binary indicator set equal to one if contract j was a cooperative agreement and zero if a grant. Our third prediction relates to the innovative performance of research contracts. We employ five dependent variables to test this prediction. $Generates\ Patent_j$ is a binary indicator set equal to one if contract j generated at least one patent and zero otherwise. $NumPatents_j$ is a count of patents generated by contract j . $Citation\text{-}Weighted\ Patents_j$ is the sum of the patents generated by contract j and the subsequent citations to those patents. $Citations/Patent_j$ is constructed as $Citation\text{-}Weighted\ Patents_j/NumPatents_j$. $Citations/Patent/Year_j$ is constructed as $Citation\text{-}Weighted\ Patents_j/NumPatents_j$ divided by the number of years since contract j was signed.

4.2.1 | Governance

The main independent variables of interest predicting the choice of governance are $Early\text{-}Stage\ Personnel_j$, which proxies for high-uncertainty projects, and $Personnel\ Expertise\ Ratio_j$, which measures the ability of government personnel to provide effective oversight on a project. For ease of explication, we discuss these in reverse order.

$Personnel\ Expertise\ Ratio_j$ is defined as the proportion of occupational categories required to conduct contract j that are available among geographically proximate client bureau personnel. We measure this using a three-step procedure. First, for each of the 59 occupational categories identified in the CPDF handbook, we create a list of distinct terms in the constituent job titles.²⁰ Next, we search for these terms in contract j ’s project description. If a term from an occupational category is found in the project description, then contract j is coded as requiring the skills of that category. Thus, each contract is characterized as drawing on a subset of the 59 occupational skill sets, with the median contract requiring skills from eight occupational categories, the mean contract requiring skills from

¹⁹Office of Personnel Management (November 14, 2014). The Guide to Data Standards, Update 16, A159-A167.

²⁰For example, the term list for the “Medical, Hospital, Dental and Public Health Group” includes terms such as health, scienc*, medic*, physician, autopsy*, dietitian, nutritionist, diagnos*, radiolog*.

12.8 occupational categories, and the SD across all contracts at 15.8. Finally, for each contract j , we calculate the proportion of requisite categories for which the sponsoring government bureau had at least one employee within a 100-mile radius of the principal research location during the year that the contract was signed.²¹ For example, if a project description signed in 2005 contained terms that occurred in occupational categories x , y , and z , and the sponsoring government bureau had at least one employee working in each of categories x and y that year within 100 miles of the research location, then the *Personnel Expertise Ratio* for that contract would be 0.67. We predicted that contracts are more likely to include monitoring/decision-rights provisions when the government has personnel who are sufficiently expert to fulfill these duties effectively; consequently, we expect the coefficient on *Personnel Expertise Ratio* to be positive.

As noted above, we expect cooperative agreements will be favored for more uncertain projects. We test this by using the functional research categories in the CPDF data to proxy for the early-stage nature of a contract. Imagine that contract j 's project description contains words that occur only in occupational category x and contract k 's project description contains words that occur only in occupational category y . If the government employees in occupational category x are clustered in the research function, while employees in occupational category y are clustered in the development function, then contract j is likely to cover a more early-stage research project than contract k .

Given this, *Early-Stage Personnel* is defined as the proportion of a bureau's geographically proximate employees in contract j 's requisite occupational categories who are assigned to a research function. We measure this in a three-step process. We start with the list of relevant occupational categories for contract j and identify client bureau personnel in those categories within 100 miles of the location of work. We then calculate the percent of these relevant employees who are also categorized in the research function. We predicted that contracts are more likely to include monitoring/decision-rights provisions when the project entails early-stage effort; consequently, we expect the coefficient on *Early-Stage Personnel* to be positive.²²

4.2.2 | Performance

The main independent variable of interest in the innovative performance of research contracts is the contractual form: *Coop Agreement*, defined above. We predicted that, conditional on government personnel assigning contracts according to project uncertainty and presence of relevant skills, cooperative agreements should outperform grants. Therefore, we expect the coefficient on *Coop Agreement* to be positive.

Control variables. We include several additional variables to control for various project, firm, bureau, and time-based characteristics. It is possible that projects with larger budgets or more firm cofunding are more likely to fall under a particular governance form or generate patents. We therefore include *Federal Funding*, defined as the dollar amount contributed by the government to support contract j , as well as *Firm/Total Funding_j*, defined as the amount contributed by the firm divided by total funding for contract j . Larger or higher-patenting firms might be more likely to generate a patent from the contracted research and/or be differentially likely to operate under a particular

²¹We use a 100-mile radius for our core estimations because this roughly corresponds to the maximum distance that one can drive twice in one workday (outbound and return) and still have time for a half-day meeting. We replicate all estimations with alternative radii of 200, 300, 400, and 500 miles.

²²To create the Early Stage Personnel variable, we link project descriptions to occupational category descriptions, and link occupational category descriptions to the functions of government scientists. This entails two assumptions. First, we assume that missing phrase matches between project descriptions and occupations are not correlated with project uncertainty. Second, we assume that project uncertainty is correlated with identifiable words in the project description and that the distribution of government research scientists in the government reflects the degree of research in a project with that same phrase description.

contractual form. To address this, we include *Large Firm_j*, which equals one if the firm is categorized by the government as a “large for-profit enterprise” in the research contract document and zero if it is coded as a “small business enterprise.” We also include *Prior Patents_j*, defined as the number of patent applications filed by the focal firm in the year preceding the signing of contract *j*. Given the skew in *Federal Funding*, *Firm/Total Funding*, and *Prior Patents*, we standardize each variable and use the *z*-scores rather than using the raw values.

A bureau's choice of contractual form, and the performance of its contract, might be affected by the degree to which its local office is managing several concurrent contracts. For a cooperative agreement in particular, this could affect performance if government researchers with relevant expertise are unable to devote as much attention and effort to contract *j*'s research project as would be optimal. We include *Coops Within 100 Miles_j* and *Coops/Personnel Ratio_j*, defined respectively as the number of in-process cooperative agreements within a 100-mile radius of contract *j*'s principal research location and the ratio of these agreements to research personnel in the local bureau. Note that the *Coops/Personnel Ratio* variable obliquely proxies for the feasibility of coordination between government researchers and the focal firm; if the government researchers are stretched too thin, then they will not be able to effectively monitor or make decisions regarding contract *j*'s research project. Contracts undertaken in different years might have different forms and outcomes due to temporal pressures on personnel; to address this we include fiscal-year fixed effects. Finally, for roughly 31% of contracts, the project description does not yield a link to any occupational categories, which precludes identifying employees with relevant skills. In these instances, we set *Personnel Expertise Ratio* equal to 0. To separate these from the qualitatively different instances in which contract *j* is linked to occupational categories and the bureau has no relevant local personnel, we also include *No Expertise_j*, a binary variable set to one if contract *j*'s project description is not linked to any occupational categories and to zero otherwise.²³

Table 1 provides summary statistics for our sample. As noted above, 12.5% of our sample contracts generate at least one patent. Almost one-quarter of the contracts are cooperative agreements. Slightly more than one-quarter of the contracts involve a large firm. The average number of concurrent cooperative agreements managed by the local relevant bureau is nearly six, which equates to nearly 0.5 agreements per local research employee. The average cooperative agreement is five times as likely to generate a patent as the average grant, and is more likely to entail early-stage effort. The average cooperative agreement is also more likely to be performed by a large, high-patenting organization in conjunction with a bureau that has relevant local expertise. Correlation matrices for the sample are provided in Table S1 of Supporting Information Appendix S1.

4.3 | Empirical strategy

To appropriately estimate the models we employ a two-stage econometric technique that first estimates the probability of selecting into a grant or cooperative agreement governance mode and then estimates the effect of this cooperative agreement “treatment” on patenting, using information from the selection model to correct for the nonrandom nature of the treatment model. The first stage provides a test for our governance predictions while the second stage provides a test for our performance predictions.

This estimation approach requires an instrumental variable for contract form in the first-stage selection model. Our instrument is *Personnel Expertise Ratio*. As described above, when there are more government research personnel with relevant expertise in the geographic area of a research location,

²³We further address the issue of missing values in the Data section of Supporting Information Appendix S1.

TABLE 1 Descriptive statistics for variables used in primary analysis

Variable	All contracts (<i>N</i> = 4,074)				Cooperative agreements only (<i>N</i> = 916)				Grants only (<i>N</i> = 3,158)			
	Mean	SD	Min.	Max.	Mean	SD	Min.	Max.	Mean	SD	Min.	Max.
Generates patent	0.125	0.330	0	1	0.326	0.469	0	1	0.066	0.249	0	1
Number of patents generated	0.322	1.963	0	59	1.096	3.964	0	59	0.097	0.439	0	6
Citation-weighted patents	1.815	23.513	0	1,082	7.480	49.073	0	1,082	0.172	1.817	0	61
Citations/patent	0.496	3.607	0	100.333	1.793	7.050	0	100.333	0.120	1.323	0	61
Citations/patent/year	0.045	0.252	0	6.271	0.146	0.475	0	6.271	0.015	0.112	0	3.813
Coop agreement	0.225	0.418	0	1	1	0	1	1	0	0	0	0
Personnel expertise ratio	0.059	0.159	0	1	0.117	0.217	0	1	0.042	0.134	0	1
Early stage personnel	0.009	0.070	0	1	0.013	0.078	0	1	0.008	0.068	0	1
Federal funding ^a	0	1	-2.331	31.400	0.214	1.342	-0.609	21.336	-0.062	0.871	-2.331	31.400
Firm/total funding ^a	0	1	-0.354	4.101	0.870	1.546	-0.354	4.101	-0.252	0.561	-0.354	4.101
Large firm	0.268	0.443	0	1	0.750	0.433	0	1	0.128	0.334	0	1
Prior year Patents ^a	0	1	-0.316	16.323	0.483	1.392	-0.316	16.323	-0.140	0.801	-0.316	16.323
Coops/personnel ratio	0.459	1.839	0	25	1.093	2.917	0	24	0.275	1.322	0	25
Coops within 100 miles	5.703	10.553	0	68	12.503	14.771	0	65	3.730	7.945	0	68
No expertise	0.312	0.463	0	1	0.084	0.278	0	1	0.378	0.485	0	1

Note. All variables calibrated to 100-mile distance.

^a z-score standardized.

the government bureau is better able to accurately evaluate progress in the research project; allocating decision and monitoring rights to the government will provide governance benefits. Therefore, cooperative agreements will be positively correlated with the number of geographically proximate government personnel with relevant technical expertise. At the same time, the presence of local government research personnel with expertise per se is unlikely to be correlated with patenting, to the extent that firm and government scientists are effectively substitutes in production. We address the robustness of this assumption in the Robustness Checks section of Supporting Information Appendix S1.

One complication in our setting is that both selection and treatment models have dichotomous dependent variables. Consequently, several conventional two-stage approaches are inappropriate because they are only robust in linear settings (Chesher, 2010; Wooldridge, 2010). Our preferred method is the inverse probability weighting regression adjustment method (IPWRA). Although not commonly used in the strategy literature, it is a mainstay in public policy scholarship (e.g., Angrist, 1998; Angrist & Pischke, 2009). Supporting Information Appendix S1 provides a more detailed explanation of the IPWRA method, along with an assessment of its strengths and weaknesses. In short, this method is appropriate when a researcher wants to estimate treatment effects from observational data combining regression adjustment with inverse probability weighting. It is also appropriate when the choice of treatment is endogenous (e.g., whether to use a cooperative agreement or grant agreement) and there is a dichotomous variable in the second stage outcome equation (e.g., whether an agreement yields a patent). A particularly attractive feature of the IPWRA method is that it does not require an omitted instrument in the first stage (as is required in two stage least squares) because it incorporates inverse probability weighting to adjust for selection into the treatment group. Finally, it is attractive because it allows both the treatment and control groups to have their own set of second-stage coefficients, recognizing that those in each group may be differentially affected by the

TABLE 2 Two-stage IPWRA probit estimation of patent generation

	First-stage model		Second-stage models	
	Model 1		Model 2A (subsample: Coops)	Model 2B (subsample: Grants)
Personnel expertise ratio (IV)	0.651 (0.000)			
Coop agreement		0.278 (0.000)		0.083 (0.000)
Early stage personnel	1.007 (0.000)		-3.660 (0.090)	0.700 (0.069)
Federal Funding	0.038 (0.276)		0.046 (0.325)	0.061 (0.084)
Firm/Total funding	0.286 (0.000)		0.077 (0.092)	0.148 (0.033)
Prior year patents	0.005 (0.849)		-0.171 (0.003)	-0.025 (0.704)
Coops/personnel ratio	-0.003 (0.839)		-0.090 (0.049)	-0.079 (0.100)
Coops within 100 miles	0.018 (0.000)		0.005 (0.461)	0.005 (0.414)
Large firm	1.270 (0.000)		0.239 (0.121)	0.304 (0.046)
No expertise	-0.663 (0.000)		0.057 (0.840)	-0.152 (0.318)
Fiscal year fixed effects	Yes		Yes	Yes

Note. $N = 4,074$ in Model 1, 916 in Model 2A, and 3,158 in Model 2B. p -Values based on heteroskedasticity-robust standard errors are reported in parentheses. Models specified using 100-mile distance variables. Coefficients on *Coop Agreement* in Models 2A and 2B are interpreted as marginal effects; that is, conditional on being a cooperative agreement or grant, respectively.

covariates. (This is somewhat analogous to interacting the treatment with each independent variable in the second stage.)

We believe that the IPWRA approach is the most relevant and most accurate statistical approach for this research question and this data set. Nevertheless, we re-estimate the models using alternative approaches that are more common in the strategy literature, including the single stage (“naïve”) probit, the two-stage least squares linear probability model, the bivariate probit, the instrumental variables probit, and the full-information-maximum-likelihood (FIML) approach with joint normality in the error terms.²⁴ The results of these robustness checks are reported in Supporting Information Appendix S1. Virtually all results using these methods are qualitatively identical to the results discussed in the next section.

5 | EMPIRICAL RESULTS

The central results of this paper are reported in Table 2. Model 1 presents the first-stage IPWRA results that emanate from the governance predictions in the theory. Models 2A and 2B present the

²⁴All of these estimation approaches are appropriate conditional on the government funding a project. If one instead assumes that the government decides among grant, cooperative agreement, and not funding the project at all, then a multilevel treatment model would be appropriate. However, we cannot observe non-funded CFPs.

second-stage IPWRA results that arise from the contract performance predictions of the theory. The first stage of this model includes all exogenous predictors as well as the instrument, *Personnel Expertise Ratio*, which predicts selection into a cooperative agreement rather than a grant. The two second-stage models incorporate the inverse of the predicted probability of selection from the first stage as a weight in the estimation as well as regression adjustment.

We begin with the first-stage results in Model 1, for which the dependent variable is *Coop Agreement*. Of particular importance, the coefficient on *Personnel Expertise Ratio* is positive ($p = 0.000$) and of substantial magnitude. Research contracts are more likely to be organized as cooperative agreements when the sponsoring bureau has relevant skills in its geographically proximate offices. The marginal effect of a bureau having personnel in all relevant areas for a project leads to an 11% increase in the likelihood of a contract being organized as a cooperative agreement rather than as a grant. In addition to being useful for our instrument, this is also consistent with the theoretical prediction that public–private contracts will be more likely to entail ongoing public supervision when the public client possesses sufficient technical expertise to effectively exert its supervisory responsibilities.

Consistent with our prediction, the coefficient on *Early-Stage Personnel* in the first stage is positive ($p = 0.000$), indicating that earlier-stage projects are more likely to be governed by cooperative agreements. If all the personnel with expertise in areas relevant to a project's required expertise work in early-stage research positions, an agreement is 16% more likely to be organized cooperatively than if none of them does. Early-stage projects, which are typically considered to be more uncertain and hence entail more unobservable effort, are thus associated with high-monitoring, high-client-decision-rights governance.

As for the control variables, contracts involving larger firms are more likely to be organized as cooperative agreements. The amount of federal funding is not associated with governance form, but the proportion of firm funding to total funding is positively associated with cooperative agreements. Finally, projects that do not identify any areas of expertise were roughly 11% less likely to be organized as cooperative agreements, suggesting the government may prefer arms-length financial support for narrowly defined projects.

We now turn to the second-stage results, for which the dependent variable is *Generates Patent*. Model 2A presents the estimated coefficients for contracts that were governed as cooperative agreements, while Model 2B presents these for contracts that were governed as grants. Both second-stage models rely on the results from the first stage. We have converted the coefficient on the treatment variable, *Coop Agreement*, to a marginal effect. Thus, the coefficient on *Coop Agreement* in Model 2A reflects the marginal effect of being a cooperative agreement versus being governed by a grant, for those contracts that were actually governed by cooperative agreement (i.e., the “average treatment effect on the treated”), while its counterpart in Model 2B reflects the effect of being a cooperative agreement vs. a grant, for those contracts that were actually governed by grant (i.e., the “average treatment effect on the untreated”). The coefficient in Model 2A is 0.278, indicating that the average cooperative agreement in our sample was nearly 28% more likely to generate a patent than it would have been if it were organized as a grant, holding all other variables at the mean. In contrast, the coefficient in Model 2B is 0.083, indicating that the average grant in our sample would have been 8% more likely to generate a patent had it been organized as a cooperative agreement. Put differently, consistent with our prediction, cooperative agreements are associated with higher innovative output than are grants, controlling for the other independent variables; this effect is more pronounced for contracts that actually were organized as cooperative agreements than for contracts that actually were organized as grants.

In model 2A, several other variables influence the likelihood that a cooperative agreement generates a patent. Of particular note, *Coops/Personnel Ratio* is negatively related to *Generates Patent*; this is consistent with the notion that as a bureau's scientific personnel get stretched thinly, they are less able to engage in smooth coordination of effort with firm personnel, thus lowering research productivity. No other coefficients on the control variables are significant at conventional thresholds except for the coefficient on *Prior Year Patents*.

Turning to Model 2B, which focuses on contracts that were organized as grants, the coefficients on *Large Firm* and *Firm/Total Funding* are both positive ($p = 0.033$ and $p = 0.046$, respectively). This indicates that larger organizations are more likely to generate patents and, consistent with incentive theory, firms that have "skin in the game" are also more likely to generate patents.

Taken together, the above results suggest that the presence or absence of relevant expertise influences the governance of research contracts, such that cooperative agreements are substantially more likely when the sponsoring government bureau has relevant skills in geographically proximate offices. In addition, early-stage projects are more likely to be governed as cooperative agreements. In turn, cooperative agreements are more likely than grants to generate patents. Had the average cooperative agreement been governed as a grant, it would have had a 28% lower probability of generating a patent. That said, those projects organized as grants would not have enjoyed a comparable increase in probability of patent generation since they are qualitatively different than the projects organized as cooperative agreements. Had the average grant been governed as a cooperative agreement, it would have had an 8% higher probability (from a lower initial baseline) of generating a patent.²⁵

5.1 | Extensions: The magnitude of innovative performance

The above estimation focuses on a binary measure of innovative performance—whether or not a research contract generates at least one patent. It is possible that other measures of innovative output will indicate different impacts of contract structure. As noted above, we constructed alternative measures of innovative output, notably *Num Patents*, *Citation-Weighted Patents*, *Citations/Patent*, and *Citations/Patent/Year*. Table 3 presents results for IPWRA estimation of models with these four dependent variables, using linear specifications in the second-stage. The table only shows the second-stage results because the first-stage results are identical to those of the IPWRA estimation in Table 2's Model 1, by definition.

For each measure of innovative output, the coefficient on *Coop Agreement* is uniformly positive ($p = 0.000$ in all models). In all four cases, this coefficient is substantially higher for contracts that actually were governed by cooperative agreements than for contracts that were governed by grants; the coefficient ranges from roughly five times larger to as much as 15 times larger. (The largest differences occur because cooperative agreements simultaneously generate more patents and more citations/patent, affecting Models 3–8.) This consistent pattern of coefficient sign and magnitude matches the core results above. Because of the extreme skewness of the dependent variables used in Table 3, we include in Supporting Information Appendix S1 the coefficients on the second-stage *Coop Agreement* variable for identical models using logged dependent variables (addressed as item

²⁵This raises a question: If the average grant would enjoy a positive (albeit small) increase in probability of patent generation if it were organized as a cooperative agreement, then why is not it governed by a cooperative agreement? The negative coefficient on *Coops/Personnel Ratio* implies that each cooperative agreement imposes a negative externality on other geographically proximate cooperative agreements due to a congestion effect. Hence, for a wide range of values, the modest bump in research productivity from converting a focal grant to cooperative agreement will be offset by the declining productivity of nearby cooperative agreements. See Supporting Information Appendix S1 for a further discussion of this point.

TABLE 3 Second-stage IPWRA results of patent generation and quality outcomes

Variable	Number of patents generated		Citation-weighted patents		Citations/patent		Citations/patent/year	
	Model 1 Coops	Model 2 Grants	Model 3 Coops	Model 4 Grants	Model 5 Coops	Model 6 Grants	Model 7 Coops	Model 8 Grants
Coop agreement	0.666 (0.000)	0.145 (0.000)	4.516 (0.000)	0.273 (0.000)	1.560 (0.000)	0.169 (0.000)	0.121 (0.000)	0.022 (0.000)
Early stage personnel	-1.285 (0.053)	0.072 (0.554)	-6.281 (0.035)	0.098 (0.714)	-1.699 (0.034)	0.035 (0.723)	-0.160 (0.037)	0.003 (0.831)
Federal funding	0.208 (0.034)	0.101 (0.013)	0.781 (0.047)	0.031 (0.374)	0.099 (0.094)	-0.006 (0.538)	0.011 (0.166)	-0.001 (0.403)
Firm/total funding	0.060 (0.349)	0.107 (0.020)	1.184 (0.061)	0.148 (0.392)	-0.062 (0.578)	0.042 (0.486)	0.001 (0.882)	0.006 (0.471)
Prior year patents	0.174 (0.224)	-0.038 (0.308)	1.658 (0.083)	-0.167 (0.066)	0.029 (0.817)	-0.063 (0.055)	-0.004 (0.648)	-0.009 (0.049)
Coops/personnel ratio	0.042 (0.696)	-0.054 (0.012)	0.034 (0.828)	-0.050 (0.220)	-0.039 (0.072)	-0.014 (0.362)	-0.005 (0.055)	-0.002 (0.334)
Coops within 100 miles	-0.009 (0.660)	0.003 (0.348)	-0.079 (0.147)	-0.008 (0.216)	-0.008 (0.232)	-0.003 (0.213)	-0.001 (0.415)	-0.000 (0.325)
Large firm	0.609 (0.057)	0.136 (0.021)	-0.244 (0.829)	0.512 (0.046)	0.291 (0.166)	0.204 (0.075)	0.033 (0.111)	0.032 (0.027)
No expertise	0.481 (0.207)	0.031 (0.399)	-1.329 (0.287)	-0.099 (0.461)	-0.106 (0.664)	-0.069 (0.417)	0.012 (0.701)	-0.001 (0.882)
Year fixed effects	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes

Note. N = 916 in odd-numbered models and 3,158 in even-numbered models. p-Values based on heteroskedasticity-robust standard errors are reported in parentheses. Models specified using 100-mile distance variables.

seven in the robustness checks below and Table S8 in Appendix S1). The results for all eight models retain their sign ($p = 0.000$ in all models).

5.2 | Robustness checks

There are several concerns that may arise from the specification and estimation strategies we employ. They include: (a) sensitivity of results to the IPWRA approach or to the chosen geographic radius for *Personnel Expertise Ratio*; (b) omitted variable bias related to temporal variance, for example, changes to the federal budget because of the American Recovery and Reinvestment Act of 2009 (ARRA, a.k.a. the federal “stimulus package”); (c) unobserved heterogeneity at the bureau level and/or firm level; (d) the empirical appropriateness of the instrumental variable; (e) skewness of patenting and citation rates; (f) the possibility that the performance results reflect more effective coordination between government and private sector scientists when they work on cooperative agreements; and (g) unobserved heterogeneity in project quality assigned to grants and cooperative agreements. We address in detail each of these concerns in Supporting Information Appendix S1. To summarize briefly: Across numerous estimation methods and a wide range of geographic radii, the results illustrated in Table 2 remain materially unchanged. Results also remain largely the same when a dummy variable to identify ARRA contracts is included, when bureau and firm effects are included, and when logged patenting and citations are used as dependent variables. Additional tests support the

empirical validity of the instrument. Finally, after controlling for endogeneity in contract choice to the best of our ability, the results remain qualitatively unchanged.

6 | DISCUSSION AND CONCLUSION

Governments throughout the world spend tens of billions of dollars annually on contractually sourced research. Yet the challenge of public contracting for private innovation, although of substantial importance, is not well understood. In this paper, we shed light on this topic. Conventional prescriptions from the “contracting for innovation” literature do not apply straightforwardly to government contracting because of restrictions that preclude the judicious use of property rights, equity investment, royalty payments, or complex contractual provisions to align parties’ incentives in the face of unobservable effort. Put differently, public contracting for private innovation is an excellent setting in which to examine one particular mechanism to induce effort: allocation of decision rights.

We predict that public retention of decision rights is more likely to be used in the face of high project uncertainty and when the available government personnel have project-relevant technical expertise. We also predict that the use of monitoring and decision rights will positively influence the likelihood that a project results in a patented innovation. We then test these predictions with data on US federal government contracts for innovation by private firms, which are generally constrained to take one of two forms: grants, in which the government retains virtually no decision rights, and cooperative agreements, in which the government retains ongoing decision rights and attendant monitoring rights.

We find empirical support for the above predictions: cooperative agreements are more likely to be used for early-stage projects than for later-stage projects, and cooperative agreements are more likely to be used when local government personnel have relevant technical expertise. Similarly, after accounting for endogeneity in governance choice due to project uncertainty and personnel expertise, we find cooperative agreements are indeed associated with higher innovative output than are grants. We interpret these results as evidence consistent with the idea that a principal chooses to govern more-uncertain projects, where the problem of noncontractible effort is higher, by retaining more decision rights and enforcing greater monitoring over the project. However, when the principal lacks the relevant know-how to properly evaluate project progress, it is preferable to leave decisions in the hands of the agent.

These results contribute to our understanding of value creation involving public organizations. Specifically, as public actors strive to generate valuable innovations to serve government needs or more broadly enhance social welfare, this study’s insights may help these actors overcome the constraints of contractual rigidities. More generally, these results contribute to the literature on public contracting (Moszoro et al., 2016) and to the debate over management of noncontractible quality for a privatized service (Hart et al., 1997; Williamson, 1999), notably for “hybrid” public management (Ménard, 2004; Rangan, Samii, & Van Wassenhove, 2006). While prior literature has identified the importance of *incentives* for the government overseer of a hybrid (Cabral et al., 2010, 2013), this study highlights the importance of that overseer’s *ability* to evaluate the effort of the private provider.

At the broadest level, recent research on contractual governance has assessed the distinct roles of coordination mechanisms versus control mechanisms (Malhotra & Lumineau, 2011; Oxley & Wada, 2009; Ryall & Sampson, 2009). For example, Lumineau and Malhotra (2011) distinguish between contractual clauses that emphasize control and those that emphasize coordination, and find that coordination-related clauses are associated with smoother functioning of contracts in the face of inter-firm friction. In a review of this literature, Lumineau (2017, p. 1561) concludes that “a strong controlling focus may raise a constant policing...of the partner’s performance.... Such a ‘carrot-and-stick’ approach with a strict oversight may create rigidity and over-monitoring.” Our study indicates that

the relationship between rigidity and reliance on control mechanisms may also flow in the opposite direction: in institutional contexts that impose contractual rigidity, control mechanisms may be more feasible than coordination mechanisms. Moreover, although this paper has many features common to Lerner and Malmendier (2010), the results here are qualitatively different. In their paper, upon termination, the funder gets the intellectual property and the fundee gets a payment. In the public-private setting, such outcomes are prohibited; instead, the fundee retains the intellectual property and the funder pays nothing.

This study also makes two empirical contributions. First, manipulating a new data set in a novel way, we develop a measure of government personnel skills at a far more microanalytic level than has been done in the past. We can thus measure the precise level of skills or capabilities, across 59 occupational categories and 19 functional areas, possessed by the personnel of a given US government bureau at a precise geographic level such as an office location, town, or any geographic radius. To the best of our knowledge, no other measure of government capabilities exists at such a level of granularity. More generally, although the capabilities literature focuses theoretically on capabilities with specific uses, data constraints have tended to restrict empirical measurement of capabilities to features such as patenting productivity (e.g., Tortoriello, 2015) or prior experience in a particular industry (e.g., Klepper & Simons, 2000). We anticipate that our measure of government capabilities can be useful in future innovation and capabilities-based research. For example, our measures could be used in determining how firm performance and innovation are affected by co-location with government scientists and facilities; how government funding of specific innovation classes affects entrepreneurship and economic development; and how relative firm performance is affected by innovation partnerships with the government.

Second, we bring to the attention of strategic management scholars the IPWRA method, a method used quite often in the public economics and program evaluation literature. We demonstrate how this method can be usefully employed in the strategic management literature when other more traditional econometric approaches are inappropriate—in particular where there are control and treatment groups and the treatment groups are endogenously selected into the treatment. This occurs, for example, in many situations where the firms' strategic choices are dichotomous and endogenous and the researcher wishes to infer predictions about the firms' performance.

There are limitations to this study. First, we have excluded from our analysis pure outsourcing arrangements, in which the government rather than the firm owns the property rights to the innovation. It would be interesting to explore whether and how our theories and empirical methods might apply to these very common arrangements. Second, our paper measures innovative output as patents. One could imagine a situation where the government might value other outputs, such as jobs, regional economic development, or a variety of political criteria. Those aspects of a potential government utility function are outside the scope of this paper. Third, we take contractual form as endogenously determined by the government, but without the influence of firms who did not win. To the extent that nonwinning applicants for innovation contracts with the government influence contract form, our analysis will not capture that influence. Finally, while our theoretical framework aspires to be universal, the empirical work is particular to the institutional details and structure of government contracting for innovation in the United States. We believe that exploring applications of the theory to other countries would be a fruitful avenue for research.

There exist a number of further unexplored questions in this vein. Does government funding enable a firm to deepen its current expertise, or to broaden its technological portfolio and capabilities? How do the human-capital capabilities of the government affect the effectiveness of private

sector research beyond patents? These questions are part of a vibrant avenue for future research on value creation and appropriation at the nexus of the government and private-firm R&D.

ACKNOWLEDGEMENTS

We are grateful to Jean-Etienne de Bettignies, J.P. Eggers, Robert Gibbons, Ricard Gil, Mitch Hoffman, Chris McKenna, Marian Moszoro, Arti Rai, Pablo Spiller, Jesper Sørensen, Giorgio Zanarone, Rosemarie Ziedonis; seminar participants at CUNEF, George Mason University, MIT, NYU, Queens University, Stanford University, the University of California, Berkeley, the University of Pompeu Fabra; and conference attendees at the annual meetings of the Society for Institutional & Organizational Economics, the Academy of Management, and the Wharton Technology & Innovation Conference, for comments on previous drafts of this paper. This study is based upon work supported by the National Science Foundation under Grants Numbers 1061600 and 1443014.

ORCID

Joshua R. Bruce  <https://orcid.org/0000-0002-5074-5276>

REFERENCES

- Aghion, P., & Tirole, J. (1994). The management of innovation. *Quarterly Journal of Economics*, 109(4), 1185–1209.
- Anderson, S. W., & Dekker, H. C. (2005). Management control for market transactions: The relation between transaction characteristics, incomplete contract design, and subsequent performance. *Management Science*, 51(12), 1734–1752.
- Angrist, J. D. (1998). Estimating the labor market impact of voluntary military service using social security data on military applicants. *Econometrica*, 66(2), 249–288.
- Angrist, J. D., & Pischke, J.-S. (2009). *Mostly harmless econometrics: An Empiricist's companion*. Princeton, NJ: Princeton University Press.
- Argyres, N., & Mayer, K. J. (2007). Contract design as a firm capability: An integration of learning and transaction cost perspectives. *Academy of Management Review*, 32(4), 1060–1077.
- Argyres, N. S., Bercovitz, J., & Mayer, K. J. (2007). Complementarity and evolution of contractual provisions: An empirical study of IT services contracts. *Organization Science*, 18(1), 3–19.
- Arora, A., & Gambardella, A. (2010). Ideas for rent: An overview of markets for technology. *Industrial and Corporate Change*, 19(3), 775–803.
- Arruñada, B., Garicano, L., & Vázquez, L. (2001). Contractual allocation of decision rights and incentives: The case of automobile distribution. *Journal of Law, Economics, and Organization*, 17(1), 257–284.
- Athey, S., & Roberts, J. (2001). Organizational design: Decision rights and incentive contracts. *American Economic Review*, 91(2), 200–205.
- Boyne, G. A. (2002). Public and private management: What's the difference? *Journal of Management Studies*, 39(1), 97–122.
- Cabral, S. (2017). Reconciling conflicting policy objectives in public contracting: The enabling role of capabilities. *Journal of Management Studies*, 54(6), 823–853.
- Cabral, S., Lazzarini, S. G., & de Azevedo, P. F. (2010). Private operation with public supervision: Evidence of hybrid modes of governance in prisons. *Public Choice*, 145(1/2), 281–293.
- Cabral, S., Lazzarini, S. G., & de Azevedo, P. F. (2013). Private entrepreneurs in public services: A longitudinal examination of outsourcing and statization of prisons. *Strategic Entrepreneurship Journal*, 7(1), 6–25.
- Cassiman, B., & Veugelers, R. (2006). In search of complementarity in innovation strategy: Internal R&D and external knowledge acquisition. *Management Science*, 52(1), 68–82.
- Chesher, A. (2010). Instrumental variable models for discrete outcomes. *Econometrica*, 78(2), 575–601.
- Cohen, W. M., & Levinthal, D. A. (1990). Absorptive capacity: A new perspective on learning and innovation. *Administrative Science Quarterly*, 35(1), 128–152.
- Decarolis, F., Giuffrida, L. M., Iossa, E., Mollisi, V., & Spagnolo, G. (2018). *Bureaucratic competence and procurement outcomes* (Working Paper No. 24201). National Bureau of Economic Research. Retrieved from <http://www.nber.org/papers/w24201>.
- Grossman, S. J., & Hart, O. D. (1986). The costs and benefits of ownership: A theory of vertical and lateral integration. *Journal of Political Economy*, 94(4), 691–719.
- Hart, O., Shleifer, A., & Vishny, R. W. (1997). The proper scope of government: Theory and an application to prisons. *The Quarterly Journal of Economics*, 112(4), 1127–1161.

- Hegde, D. (2014). Tacit knowledge and the structure of license contracts: Evidence from the biomedical industry. *Journal of Economics and Management Strategy*, 23(3), 568–600.
- Hourihane, M., & Parkes, D. (2015). Federal R&D in the FY 2016 budget: An overview. *American Association for the Advancement of Science* Retrieved from <https://www.aasas.org/fy16budget/federal-rd-fy-2016-budget-overview>
- Kivleniece, I., & Quelin, B. V. (2012). Creating and capturing value in public-private ties: A private actor's perspective. *Academy of Management Review*, 37(2), 272–299.
- Klein, P. G., Mahoney, J. T., McGahan, A. M., & Pitelis, C. N. (2013). Capabilities and strategic entrepreneurship in public organizations. *Strategic Entrepreneurship Journal*, 7(1), 70–91.
- Klepper, S., & Simons, K. L. (2000). The making of an oligopoly: Firm survival and technological change in the evolution of the US tire industry. *Journal of Political Economy*, 108(4), 728–760.
- Lan, Z., & Rainey, H. G. (1992). Goals, rules, and effectiveness in public, private, and hybrid organizations: More evidence on frequent assertions about differences. *Journal of Public Administration Research and Theory*, 2(1), 5–28.
- Lerner, J., & Malmendier, U. (2010). Contractibility and the design of research agreements. *American Economic Review*, 100(1), 214–246.
- Lerner, J., & Merges, R. P. (1998). The control of technology alliances: An empirical analysis of the biotechnology industry. *The Journal of Industrial Economics*, 46(2), 125–156.
- Levin, J., & Tadelis, S. (2010). Contracting for government services: Theory and evidence from US cities. *The Journal of Industrial Economics*, 58(3), 507–541.
- Lumineau, F. (2017). How contracts influence trust and distrust. *Journal of Management*, 43(5), 1553–1577.
- Lumineau, F., & Malhotra, D. (2011). Shadow of the contract: How contract structure shapes interfirm dispute resolution. *Strategic Management Journal*, 32(5), 532–555.
- Malhotra, D., & Lumineau, F. (2011). Trust and collaboration in the aftermath of conflict: The effects of contract structure. *Academy of Management Journal*, 54(5), 981–998.
- Mayer, K. J., & Salomon, R. M. (2006). Capabilities, contractual hazards, and governance: Integrating resource-based and transaction cost perspectives. *Academy of Management Journal*, 49(5), 942–959.
- McKenna, C. D. (2006). *The World's newest profession: Management consulting in the twentieth century*. Cambridge, England: Cambridge University Press.
- Ménard, C. (2004). The economics of hybrid organizations. *Journal of Institutional and Theoretical Economics*, 160(3), 345–376.
- Miller, A. S. (1955). Government contracts and social control: A preliminary inquiry. *Virginia Law Review*, 41(1), 27–58.
- Moe, T. M. (1989). The politics of bureaucratic structure. In J. E. Chubb & P. E. Peterson (Eds.), *Can the government govern?* (pp. 267–329). Washington, D.C.: Brookings Institution.
- Moe, T. M. (1990). Political institutions: The neglected side of the story. *Journal of Law, Economics, and Organization*, 6, 213–253.
- Moszoro, M., Spiller, P. T., & Stolorz, S. (2016). Rigidity of public contracts. *Journal of Empirical Legal Studies*, 13(3), 396–427.
- National Audit Office. (2013). *Research and Development funding for science and technology in the UK*. London, England: Memorandum for the House of Commons Science and Technology Committee.
- Oxley, J. E. (1997). Appropriability hazards and governance in strategic alliances: A transaction cost approach. *Journal of Law, Economics, and Organization*, 13(2), 387–409.
- Oxley, J. E., & Wada, T. (2009). Alliance structure and the scope of knowledge transfer: Evidence from US-Japan agreements. *Management Science*, 55(4), 635–649.
- Pisano, G. P. (1990). The R&D boundaries of the firm: An empirical analysis. *Administrative Science Quarterly*, 35(1), 153–176.
- Quelin, B. V., Kivleniece, I., & Lazzarini, S. (2017). Public-private collaboration, hybridity and social value: Towards new theoretical perspectives. *Journal of Management Studies*, 54, 763–792.
- Rangan, S., Samii, R., & Van Wassenhove, L. N. (2006). Constructive partnerships: When alliances between private firms and public actors can enable creative strategies. *Academy of Management Review*, 31(3), 738–751.
- Reuer, J. J., & Ariño, A. (2007). Strategic alliance contracts: Dimensions and determinants of contractual complexity. *Strategic Management Journal*, 28(3), 313–330.
- Reuer, J. J., Ariño, A., & Mellewigt, T. (2006). Entrepreneurial alliances as contractual forms. *Journal of Business Venturing*, 21(3), 306–325.
- Reuer, J. J., & Devarakonda, S. V. (2016). Mechanisms of hybrid governance: Administrative committees in non-equity alliances. *Academy of Management Journal*, 59(2), 510–533.
- Ryall, M. D., & Sampson, R. C. (2009). Formal contracts in the presence of relational enforcement mechanisms: Evidence from technology development projects. *Management Science*, 55(6), 906–925.
- Sampson, R. C. (2004a). Organizational choice in R&D alliances: Knowledge-based and transaction cost perspectives. *Managerial and Decision Economics*, 25(6/7), 421–436.
- Sampson, R. C. (2004b). The cost of misaligned governance in R&D alliances. *Journal of Law, Economics, and Organization*, 20(2), 484–526.
- Spiller, P. T., & Moszoro, M. (2014). Third-party opportunism and the theory of public contracts: Operationalization and applications. In E. Rousseau & J.-M. Glachant (Eds.), *The manufacturing of markets: Legal, political and economic dynamics* (pp. 229–252). New York, NY: Cambridge University Press.

- Teece, D. J. (1986). Profiting from technological innovation: Implications for integration, collaboration, licensing and public policy. *Research Policy*, 15(6), 285–305.
- Tortoriello, M. (2015). The social underpinnings of absorptive capacity: The moderating effects of structural holes on innovation generation based on external knowledge. *Strategic Management Journal*, 36(4), 586–597.
- Williamson, O. E. (1999). Public and private bureaucracies: A transaction cost economics perspectives. *Journal of Law, Economics, and Organization*, 15(1), 306–342.
- Wooldridge, J. M. (2010). *Econometric analysis of cross section and panel Data* (2nd ed.). Cambridge, MA: MIT Press.
- Xiao, W., & Xu, Y. (2012). The impact of royalty contract revision in a multistage strategic R&D alliance. *Management Science*, 58(12), 2251–2271.

SUPPORTING INFORMATION

Additional supporting information may be found online in the Supporting Information section at the end of the article.

How to cite this article: Bruce JR, de Figueiredo JM, Silverman BS. Public contracting for private innovation: Government capabilities, decision rights, and performance outcomes. *Strat Mgmt J*. 2019;40:533–555. <https://doi.org/10.1002/smj.2973>